

# Design Document: Memory Management Simulator

System Architecture & Implementation Report

January 8, 2026

## 1 Introduction

This document describes the design and implementation of a comprehensive Memory Management Simulator. The system models the interactions between an Operating System's memory management unit (MMU), a multilevel CPU cache hierarchy, and physical memory allocation. The goal is to simulate OS-level behavior including dynamic memory allocation, cache replacement policies, and virtual memory paging using user-space abstractions.

## 2 System Architecture

The simulator is modular, consisting of three primary subsystems that interact hierarchically:

1. **Memory Management Unit (MMU):** Handles virtual-to-physical address translation and page fault management.
2. **Cache Controller:** Manages a two-level (L1/L2) cache hierarchy to simulate memory access latency reduction.
3. **Physical Memory Allocator:** Manages the raw physical memory space using dynamic allocation strategies.

### 2.1 Data Flow

The simulation follows a strictly defined order of operations for memory access:

User Request(Virtual/Physical Addr) → MMU → Physical Addr → Cache Controller  
→ Physical Memory

## 3 Component Design

### 3.1 1. Physical Memory Allocation

The physical memory is simulated as a contiguous block of bytes. The simulator implements a **Standard Allocator** using a linked-list data structure to track memory usage.

- **Data Structure:** A singly linked list of **Block** structures. Each block contains metadata: `start_address`, `size`, `is_free`, and a reference to the next block.
- **Allocation Strategies:** The system supports three configurable strategies:
  - **First Fit:** Allocates the first free block that satisfies the request.
  - **Best Fit:** Scans all blocks to find the smallest free block that fits the request, minimizing wasted space.

- **Worst Fit:** Allocates the largest available free block, leaving large gaps for future use.
- **Deallocation & Coalescing:** When a block is freed (by physical address), the system marks it as free and immediately iterates through the list to merge adjacent free blocks to reduce external fragmentation.

### 3.2 2. Cache Simulation

The system simulates a configurable multilevel cache (L1 and L2).

- **Configuration:** Each level supports configurable total size, block size, and associativity (Direct-Mapped or Set-Associative). Size of L2 cache is set to 8 times the size of L1 for simplicity.
- **Addressing:** Physical addresses are split into *Tag*, *Index*, and *Offset* bits based on the cache geometry.
- **Replacement Policy:** The cache implements the **Least Recently Used (LRU)** policy. A global access counter is used to track the "age" of cache lines, evicting the line with the smallest timestamp upon a conflict.
- **Write Policy:** The simulation assumes a write-allocate policy where write misses bring the block into the cache.

### 3.3 3. Virtual Memory (MMU)

The MMU provides a layer of abstraction between the user and physical memory using Paging.

- **Page Table:** Implemented as a `std::map` mapping Virtual Page Numbers (VPN) to Page Table Entries (PTE).
- **Page Table Entry (PTE):** Tracks the Valid bit, Physical Frame Number (PFN), Dirty bit, and Last Access Time.
- **Address Translation:**

$$VPN = \frac{VirtualAddress}{PageSize}, \quad Offset = VirtualAddress \pmod{PageSize}$$
- **Page Fault Handling:** If a VPN is invalid:
  1. The MMU requests a free frame from the Physical Memory Allocator.
  2. If memory is full, the **LRU Page Replacement** algorithm evicts the least recently used page.
  3. If the victim page is "Dirty", a disk write is simulated (symbolic logging).
  4. The new mapping is updated in the Page Table.

## 4 Implementation Details

The project is implemented in C++ using a modular class structure:

- **MemorySimulator:** Manages the linked list and physical byte array.
- **CacheLevel & CacheController:** Manage cache logic and L1/L2 coordination.
- **MMU:** Manages translation and orchestrates interactions between Cache and Memory.
- **main.cpp:** Provides a Command Line Interface (CLI) for user interaction.

## 5 Assumptions and Limitations

- **Memory Units:** All memory sizes and addresses are represented in Bytes.
- **Page Size:** The Page Size is assumed to be equal to the Frame Size.
- **Disk Backing:** Disk storage is simulated symbolically; evicted pages are effectively destroyed unless re-loaded, and "loading" simply allocates a new zeroed frame.
- **Single Process:** The simulation models a single process environment; context switching is not supported.